

# M2-BIG DATA

## High Performance Computing Profiling

### Lab 1



Réalisé par :  
Gaby Maroun

Encadré par :  
Dr. Etancelin JM

16 octobre 2020

Explore 3 ways of profiling a code on the Bilbo cluster :

## 1 How to use the cluster :

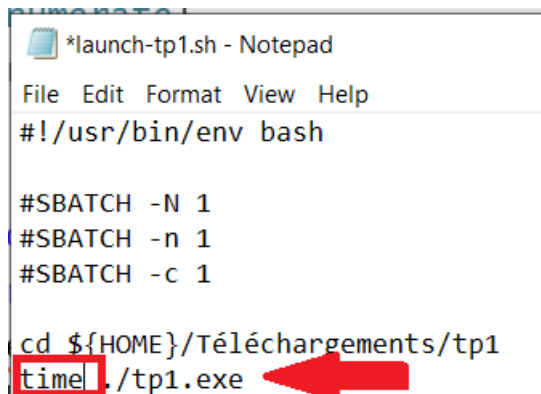
1. Login to cluster (x2go or ssh)
2. Prepare the codes, compile, explore results,
3. Submit jobs on to the resource manager (SLURM).

See more detailed information on the course page.

## 2 System profiling

Get the sources of the program from the course page. It is better to copy the archive file and extract on the cluster `tar xvf tp1.tar.gz`

1. Build the code : **cd tp1** then **make**
2. Launch the code on computing resources : **sbatch launch-tp1.sh** The file **launch-tp1.sh** is a submission script containing resources need description and commands to launch on ressources. For example :



```
*launch-tp1.sh - Notepad
File Edit Format View Help
#!/usr/bin/env bash

#SBATCH -N 1
#SBATCH -n 1
#SBATCH -c 1

cd ${HOME}/Téléchargements/tp1
time ./tp1.exe
```

3. Adapt your submission script in order to run the program through the system command time. Explore results and output in the file slurm-NNNNN.out and comment your results.

```
slurm-7999.out - Notepad
File Edit Format View Help
TP1
Read 4779 Nodes. Read 9558 elements (total)
Assemble matrices with 4779 compute points using gauss3pt2d integration formula on 9379 elements.
Writing file tp1_000010.h5
Writing file tp1_000020.h5
Writing file tp1_000030.h5
Writing file tp1_000040.h5
Writing file tp1_000050.h5
Writing file tp1_000060.h5
Writing file tp1_000070.h5
Writing file tp1_000080.h5
Writing file tp1_000090.h5
Writing file tp1_000100.h5
real 0m10,742s
user 0m9,820s
sys 0m0,068s
```

In the **BLUE** part, we can see the amount of nodes and elements read, and the amount of compute points using gauss3ptd2d integration formula to assemble matrices on the 9379 elements.

In the **ORANGE** part, a list containing the written in files ".h5".

In the **YELLOW** part, we see the result of the **time** command represented by its three parameters real, user and sys. As of the precedent results, we can assume that the system or CPU took 0,068s to execute the code from the c file while in reality it took 10,742s.

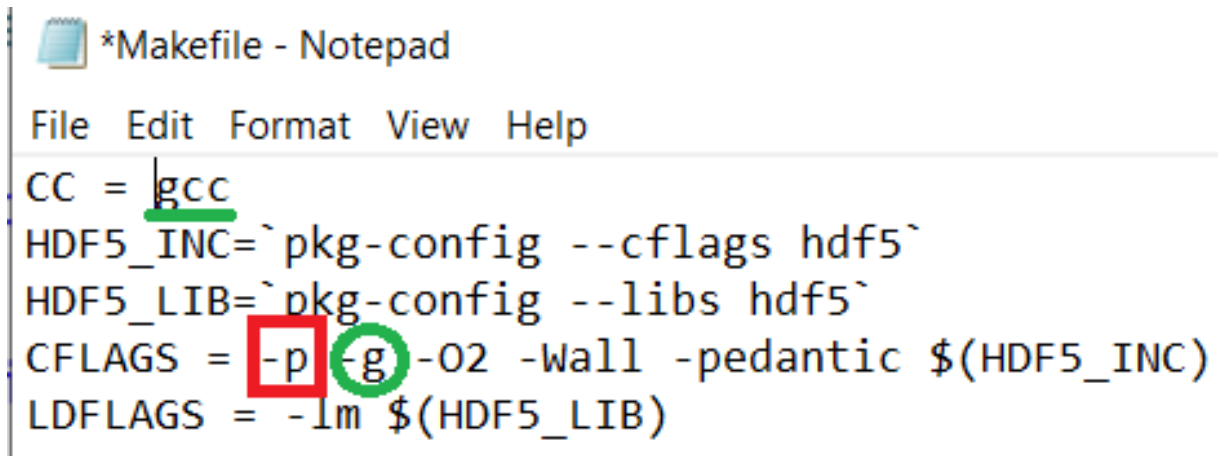
### 3 Gprof profiling

1. Clean the sources before re-compiling : **make clean**

```
gmaroun@slurm-ens-frontal:~/Téléchargements/tp1$ make clean
rm -f iterative.o mesh2dp1.o elem2dp1.o sparsematrix.o hdfio.o vector.o direct.o tp1.o matrix.o elem1dp1.o tp1.exe
```

After the cleaning, the `tp1.exe` won't be found anymore so we use the **make** to recreate it.

2. Recompile the code using the **gcc** profiling option **-p** (check that the **-g** option is also present) by modifying the CFLAGS variable in the **Makefile**



```
*Makefile - Notepad
File Edit Format View Help
CC = gcc
HDF5_INC=`pkg-config --cflags hdf5`
HDF5_LIB=`pkg-config --libs hdf5`
CFLAGS = -p -g -O2 -Wall -pedantic $(HDF5_INC)
LDFLAGS = -lm $(HDF5_LIB)
```

3. Re-launch your execution script without the use of time command and explore the output.  
Check that the file **gmon.out** has been created properly

```
gmaroun@slurm-ens-frontal:~/Téléchargements/tp1$ cat slurm-8056.out
gmaroun@slurm-ens-frontal:~/Téléchargements/tp1$ gedit slurm-8056.out
gmaroun@slurm-ens-frontal:~/Téléchargements/tp1$ ls
direct.c  elem1dp1.h  elem2dp1.o  hdfio.o  iterative.o  matrix.h  mesh2dp1.o  slurm-8001.out  slurm-8041.out  sparsematrix.h  tp1_000030.h5  tp1_000070.h5  tp1.c  tp1.xmf
direct.h  elem1dp1.o  gmon.out    launch-tp1.sh  matrix.o  meshtp10_0.05-0.01.msh  slurm-8004.out  slurm-8050.out  sparsematrix.o  tp1_000040.h5  tp1_000080.h5  tp1.exe  vector.c
direct.o  elem2dp1.c  hdfio.c    iterative.c  Makefile  mesh2dp1.c  slurm-7975.out  slurm-8009.out  slurm-8056.out  tp1_000010.h5  tp1_000050.h5  tp1_000090.h5  tp1_mesh.h5  vector.h
elem1dp1.c  elem2dp1.h  hdfio.h    iterative.h  matrix.c  mesh2dp1.h  slurm-7999.out  slurm-8037.out  sparsematrix.c  tp1_000020.h5  tp1_000060.h5  tp1_000100.h5  tp1.o  vector.o
```

#### 4. Explore profile using `gprof -flat-profile ./tp1.exe gmon.out`

```
gmaroun@slurm-ens-frontal:~/Téléchargements/tp1$ cat slurm-8086.out
Flat profile:
```

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self us/call	total us/call	name
55.99	5.85	5.85				<b>solveSparseGaussSeidel</b>
35.03	9.51	3.66	74857	48.90	48.90	prodSMV
6.03	10.14	0.63	74857	8.42	8.42	norme2
2.78	10.43	0.29	74857	3.87	3.87	axpy
0.10	10.44	0.01				createMesh2dp1FromFile
0.10	10.45	0.01				initSMMS
0.00	10.45	0.00	337644	0.00	0.00	prodMV
0.00	10.45	0.00	255521	0.00	0.00	idxOfNbor2dp1
0.00	10.45	0.00	168822	0.00	0.00	elem2dp1_phihat_0
0.00	10.45	0.00	168822	0.00	0.00	elem2dp1_phihat_1
0.00	10.45	0.00	168822	0.00	0.00	elem2dp1_phihat_2
0.00	10.45	0.00	112548	0.00	0.00	elem2dp1_dxphi_hat_0
0.00	10.45	0.00	112548	0.00	0.00	elem2dp1_dxphi_hat_1
0.00	10.45	0.00	112548	0.00	0.00	elem2dp1_dxphi_hat_2
0.00	10.45	0.00	112548	0.00	0.00	elem2dp1_dyphi_hat_0
0.00	10.45	0.00	112548	0.00	0.00	elem2dp1_dyphi_hat_1
0.00	10.45	0.00	112548	0.00	0.00	elem2dp1_dyphi_hat_2
0.00	10.45	0.00	84411	0.00	0.00	insertNborIdx
0.00	10.45	0.00	48465	0.00	0.00	elem1dp1_phihat_0
0.00	10.45	0.00	48465	0.00	0.00	elem1dp1_phihat_1
0.00	10.45	0.00	48249	0.00	0.00	computePtFromRefPt2dp1_edge
0.00	10.45	0.00	18758	0.00	0.00	createMatrix
0.00	10.45	0.00	18758	0.00	0.00	destroyMatrix
0.00	10.45	0.00	9379	0.00	0.00	createElem2dp1
0.00	10.45	0.00	9379	0.00	0.00	destroyElem2dp1
0.00	10.45	0.00	206	0.00	0.00	createVector
0.00	10.45	0.00	205	0.00	0.00	destroyVector
0.00	10.45	0.00	179	0.00	0.00	createElem2dp1_edge
0.00	10.45	0.00	10	0.00	0.00	stepXDMFNScalar
0.00	10.45	0.00	10	0.00	0.00	stepXDMFTimedGridScalar
0.00	10.45	0.00	10	0.00	0.00	writeHDF5AnyData
0.00	10.45	0.00	10	0.00	0.00	writeXDMFMesh2DPk
0.00	10.45	0.00	4	0.00	0.00	createSMatrix
0.00	10.45	0.00	2	0.00	0.00	initSMatrixStructureFromMesh2dp1
0.00	10.45	0.00	1	0.00	0.00	closeXMF
0.00	10.45	0.00	1	0.00	0.00	openXMF
0.00	10.45	0.00	1	0.00	0.00	writeHDF5Mesh2dp1

%  
time      the percentage of the total running time of the  
          program used by this function.

%  
time      the percentage of the total running time of the  
          program used by this function.

cumulative  
seconds    a running sum of the number of seconds accounted  
          for by this function and those listed above it.

self  
seconds    the number of seconds accounted for by this  
          function alone. This is the major sort for this  
          listing.

calls      the number of times this function was invoked, if  
          this function is profiled, else blank.

self  
ms/call    the average number of milliseconds spent in this  
          function per call, if this function is profiled,  
          else blank.

total  
ms/call    the average number of milliseconds spent in this  
          function and its descendents per call, if this  
          function is profiled, else blank.

name       the name of the function. This is the minor sort  
          for this listing. The index shows the location of  
          the function in the gprof listing. If the index is  
          in parenthesis it shows where it would appear in  
          the gprof listing if it were to be printed.

Copyright (C) 2012-2018 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification,  
are permitted in any medium without royalty provided the copyright  
notice and this notice are preserved.

5. Find the most time-consuming function. Comment and justify your answer.

*As we can see from the screenshot that the most time-consuming function is the "solveSparseGaussSeidel" taking an all 5,85sec that is 55,99% of the total amount. So the function is taking an amount of time more than the half of the total which is quite a lot regarding the tasks. It's in these situation where we start thinking of parallelization as a solution.*

## 4 Scalasca/ScoreP

Profiling by automatic code instrumentation.

1. Recompile with **scorep** : **make clean** and then change **CC="scorep gcc"** and remove the **-p** added in previous part in the **Makefile**. In order to activate **scalasca**, you must load the module : **module load scalasca** in the terminal for re-building.



\*Makefile - Notepad

File Edit Format View Help

```
CC = scorep gcc
HDF5_INC=`pkg-config --cflags hdf5`
HDF5_LIB=`pkg-config --libs hdf5`
CFLAGS = -g -O2 -Wall -pedantic $(HDF5_INC)
LDFLAGS = -lm $(HDF5_LIB)
```

2. Modify the submission script to launch the code through scalasca program : add the module activation before running the program with **scalasca -analyze ./tp1.exe**



\*launch-tp1.sh - Notepad

File Edit Format View Help

```
#!/usr/bin/env bash
```

```
#SBATCH -N 1
```

```
#SBATCH -n 1
```

```
#SBATCH -c 1
```

```
cd ${HOME}/Téléchargements/tp1
```

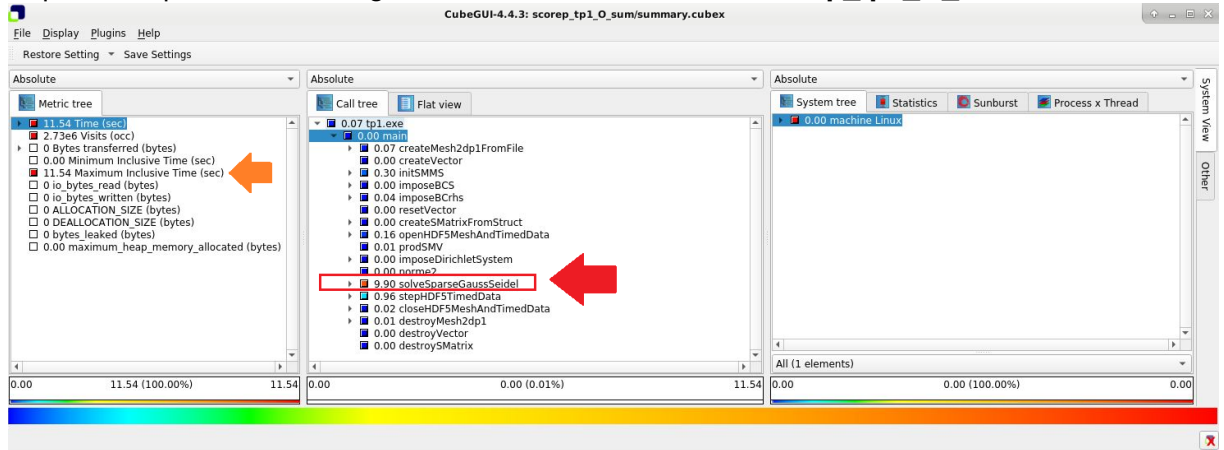
```
scalasca -analyze ./tp1.exe
```



- Check that the folder **scorep\_tp1\_O\_sum** has been created properly.

```
gmaroun@slurm-ens-frontal:~/Téléchargements/tp1$ gedit slurm-8106.out
gmaroun@slurm-ens-frontal:~/Téléchargements/tp1$ ls
direct.c  elem1dp1.o  hdfio.c      iterative.h  matrix.h      mesh2dp1.o  0.65-0.01.msh  slurm-8084.out  slurm-8086.out  sparsematrix.o  tp1_000050.h5  tp1_000100.h5  tp1.xmf
direct.h  elem2dp1.c  hdfio.h      iterative.o  matrix.o      mesh2dp1.c  scorep_tp1_O_sum  slurm-8089.out  slurm-8086.out  tp1_000010.h5  tp1_000060.h5  tp1.c          vector.c
direct.o  elem2dp1.h  hdfio.o      launch-tp1.sh  mesh2dp1.h  mesh2dp1.o  slurm-7999.out  slurm-8037.out  slurm-8106.out  tp1_000020.h5  tp1_000070.h5  tp1.exe        vector.h
elem1dp1.c  elem2dp1.o  integration.h  Makefile     mesh2dp1.h  slurm-7999.out  slurm-8041.out  sparsematrix.c  tp1_000030.h5  tp1_000080.h5  tp1_mesh.h5    vector.o
elem1dp1.h  gmon.out   iterative.c   matrix.c     mesh2dp1.o  slurm-8001.out  slurm-8050.out  sparsematrix.h  tp1_000040.h5  tp1_000090.h5  tp1.o
```

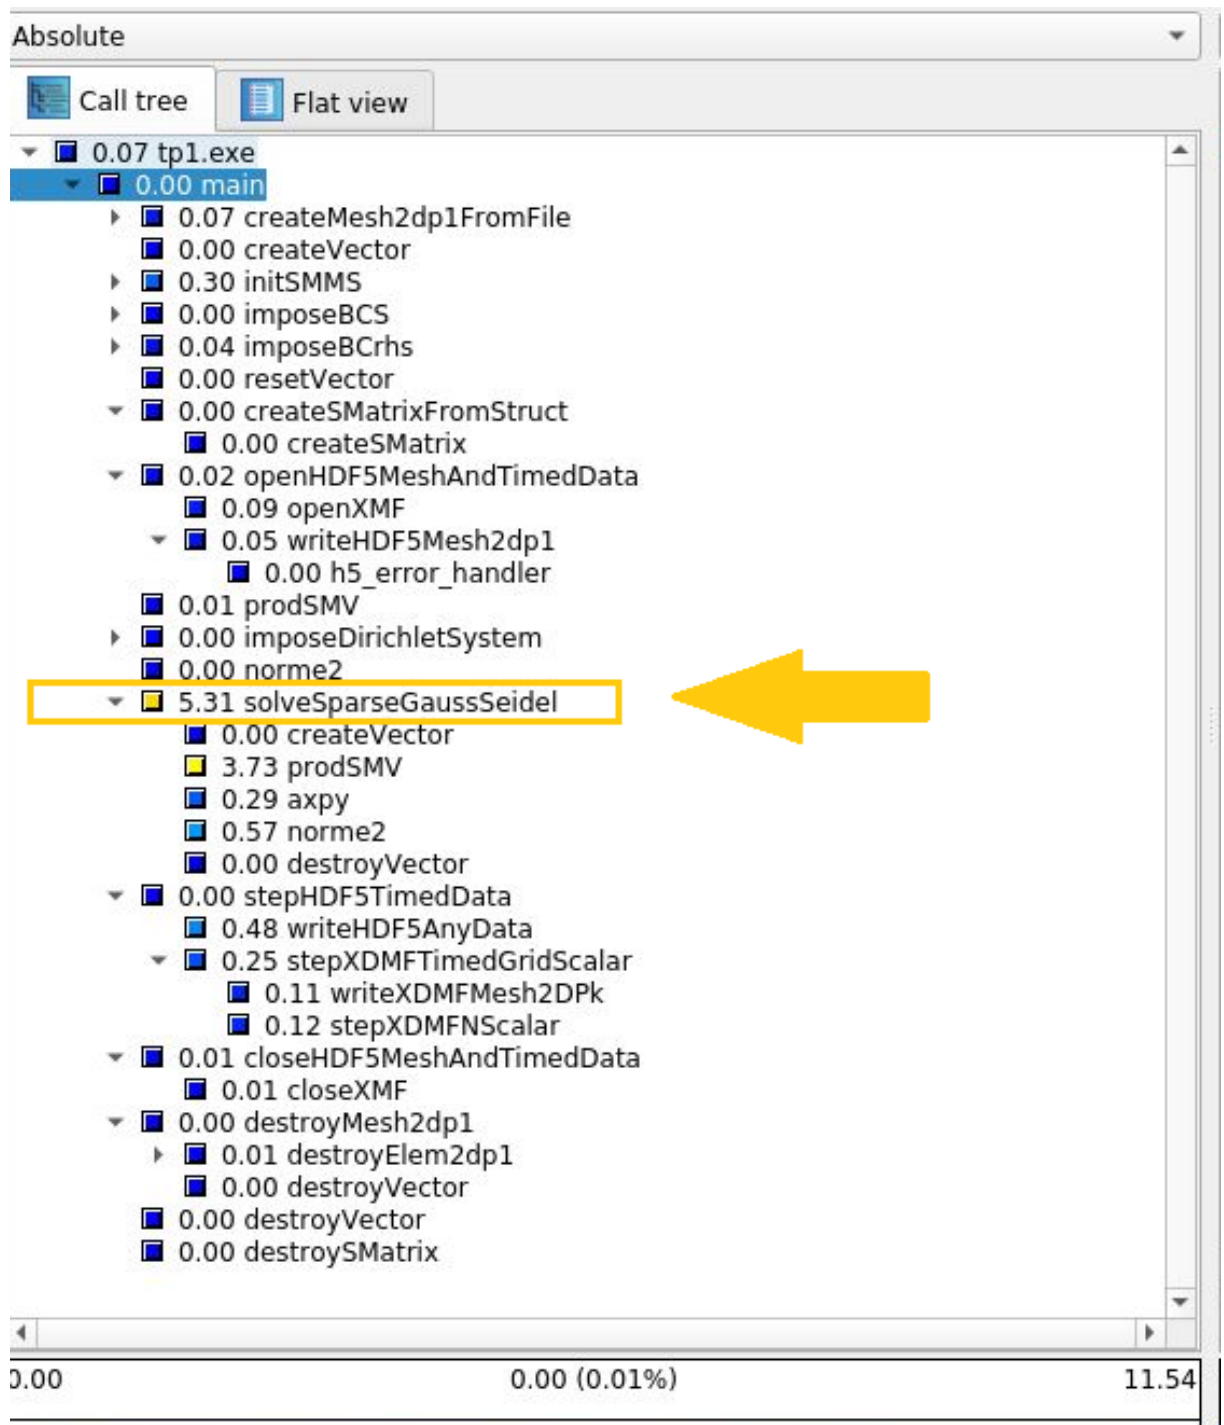
- Explore the profile data using GUI : **scalasca -examine ./scorep\_tp1\_O\_sum**



By using the **CubeGUI** with **scalasca**, we can be able to see the time of consumption taken by each function when in parallel mode.



5. Find the most time-consuming function. Comment and justify your answer. Is it in coherence with the results of the previous part ?



The most time-consuming function is also the "solveSparseGaussSeidel" same as before, so yes, in coherence. However, the time consumed while on parallel is 5,31sec, which is less by more than half of the second than the previous result.

6. Explain what function or what part of the code should be parallelized. What would be the ideal speedup based on these sequential results.

*As it has already been understood, the “solveSparseGaussSeidel” function takes the most consumption time of the total and that’s what makes it perfectly fit to be parallelized.*

*The speedup based on these sequential results is  $S_p = T_1/T_p = 5.85/5.31 = 1.101sec$ .*

*With an efficiency of  $E_p = S_p/p = 1.101sec/6 = 0.1835 < 1$ . We can conclude that it isn’t so much efficient while in parallel mode.*



*As we can see from this photo, there are 6 processors working on this PC which means the ideal and most efficient speedup would be to have a  $T_P = T_1/p = 5.85/6 = 0.975sec$*

7. Performs a second analysis after changing the variable **outputfreq** to 0 in the **tp1.c** file in order to deactivate writing the results in output files. Same question as previous one with this new experiment. Compare and comment the results.

*We notice that the time consumed by the “prodSMV” function has dropped by 0.21s, and “solveSparseGaussSeidel” function still takes a lot of time to get executed with the parallelization as a solution*

La fin.